

Jeetendra Kumar Patel

AI Engineer

+91 7000957216 | jeetendrapatel1711@gmail.com | Noida, Uttar Pradesh | linkedin.com/in/jeetendra-kumar-patel-b73650239 | jeetendra.vercel.app

PROFESSIONAL SUMMARY

Associate Software Engineer at Venera Technologies and primary engineer on the company's first generative-AI feature - a RAG-based chatbot on Qwen models that generates QC templates for Pulsar, their legacy on-premise product, from plain language. Built it from scratch with an external AI consultant and own both the implementation and the testing. Before that, a 6-month internship that delivered 150 automated Python test scripts, 120 validated API endpoints and zero critical defects shipped across AWS cloud-native microservices.

TECHNICAL SKILLS

AI & LLM Engineering: RAG, Qwen (open-weight), LLMs (OpenAI API), LangChain, LangGraph, Hugging Face Transformers, Fine-tuning (LoRA), Pinecone, Embeddings, Prompt Engineering, Structured Output, LLM Evaluation, NLP, Streamlit

Languages: Python, SQL / MySQL, Java, Kotlin, C / C++

Testing & Quality: pytest, Selenium WebDriver, Page Object Model, Selenium Grid, Postman, JUnit, SDLC, BVA

Cloud & DevOps: AWS, Azure, Jenkins, CI/CD, Git / GitHub, Bitbucket, Gradle, Maven, Zoho

EXPERIENCE

Associate Software Engineer | Aug 2026 - Present

Venera Technologies, Noida

- Primary engineer on Venera's first generative-AI feature: a retrieval-augmented chatbot that generates QC templates for Pulsar, the legacy on-premise product, from a plain-language description instead of manual configuration. Built from scratch and owned end to end since August.
- Built on Qwen open-weight models, working alongside an external professional AI consultant - chosen so the model runs within the on-premise product rather than calling a third-party API, which the legacy deployment and its customers require.
- Own the implementation: retrieval over the existing template corpus and schema so generations are grounded in real templates, prompt design, and structured-output validation so a malformed or invalid template can never reach the product.
- Own the testing too - the evaluation set the chatbot is scored against, covering both the templates it must generate correctly and the requests it must refuse rather than guess at.
- Continue to contribute across the full SDLC using OOP principles and QA best practices, and work with dev and product to define acceptance criteria and AWS cloud-native release readiness.

Graduate Trainee / QA Engineer | Feb 2026 - Aug 2026

Venera Technologies, Noida

- Engineered 150 automated Python test scripts (Selenium + POM) - cut manual regression time by 65%, reached 78% P1/P2 coverage and caught 3 critical regressions pre-release.
- Validated 120 REST API endpoints via Postman across AWS microservices (75 on Quasar, the cloud QC product; 45 on Pulsar, the on-premise one); 98.3% pass rate and 2 production-blocking defects resolved before deployment.
- Delivered zero critical defects to production across 7 builds; intercepted a severity-1 503 failure and coordinated a same-day hotfix, saving roughly 4 hours of user-facing downtime.
- Managed 79 defects in Zoho (61 resolved, 10 critical) and maintained test assets in Bitbucket with structured PR reviews.

Android Developer (Internship) | Prior to Feb 2026

NullClass, Remote

- Built and optimised Android mobile applications, gaining hands-on experience in Kotlin, UI performance tuning, Firebase integration and mobile SDLC delivery.

SELECTED PROJECTS

This Portfolio + Resume Assistant | AI / Full-stack

- The site you're reading. A Flask app with a retrieval-grounded assistant that answers recruiter questions from my resume - sub-millisecond local retrieval, nine interchangeable model providers, and a fallback that still answers with zero API keys. **Impact:** 0 vendor lock-in. **Stack:** Python, Flask, BM25, RAG, Groq, Gemini, Vanilla JS

Document Q&A Chatbot (RAG) | AI / LLM

- Retrieval-augmented Q&A over a 200-page knowledge base with grounded, citation-friendly answers. **Impact:** ~70% faster retrieval. **Stack:** LangChain, OpenAI API, Pinecone, Streamlit, Python

NLP Pipelines with Hugging Face | AI / NLP

- Transformer-based text pipelines for classification and summarisation, with prompt variants scored against a labelled set instead of judged by eye. **Impact:** measured, not guessed. **Stack:** Hugging Face, Transformers, Python, LangGraph

EDUCATION

B.E. Computer Science Engineering | Graduated June 2026

Chandigarh University, Punjab | CGPA 8.4 / 10

CERTIFICATIONS & TRAINING

- The AI Engineer Bootcamp 2026** - 365 Careers, Udemy (In Progress). LLMs, NLP, LangChain, LangGraph, Hugging Face, Pinecone, OpenAI API, RAG, Streamlit
- Completed:** Android Developer Badges (Google); Database Fundamentals (LinkedIn); Design Thinking (LinkedIn)